

# Transactional outbox pattern

Helsinki Gophers meetup

22 Jan 2025

Nikolay Kuznetsov

@nikolayk812

# About me

Senior software engineer @Zalando Oy

C → Java → Kotlin → Go

Conference speaker in 2019/20

ice-skating, kayaking, hiking, chess

learning Finnish, Swedish, Italian

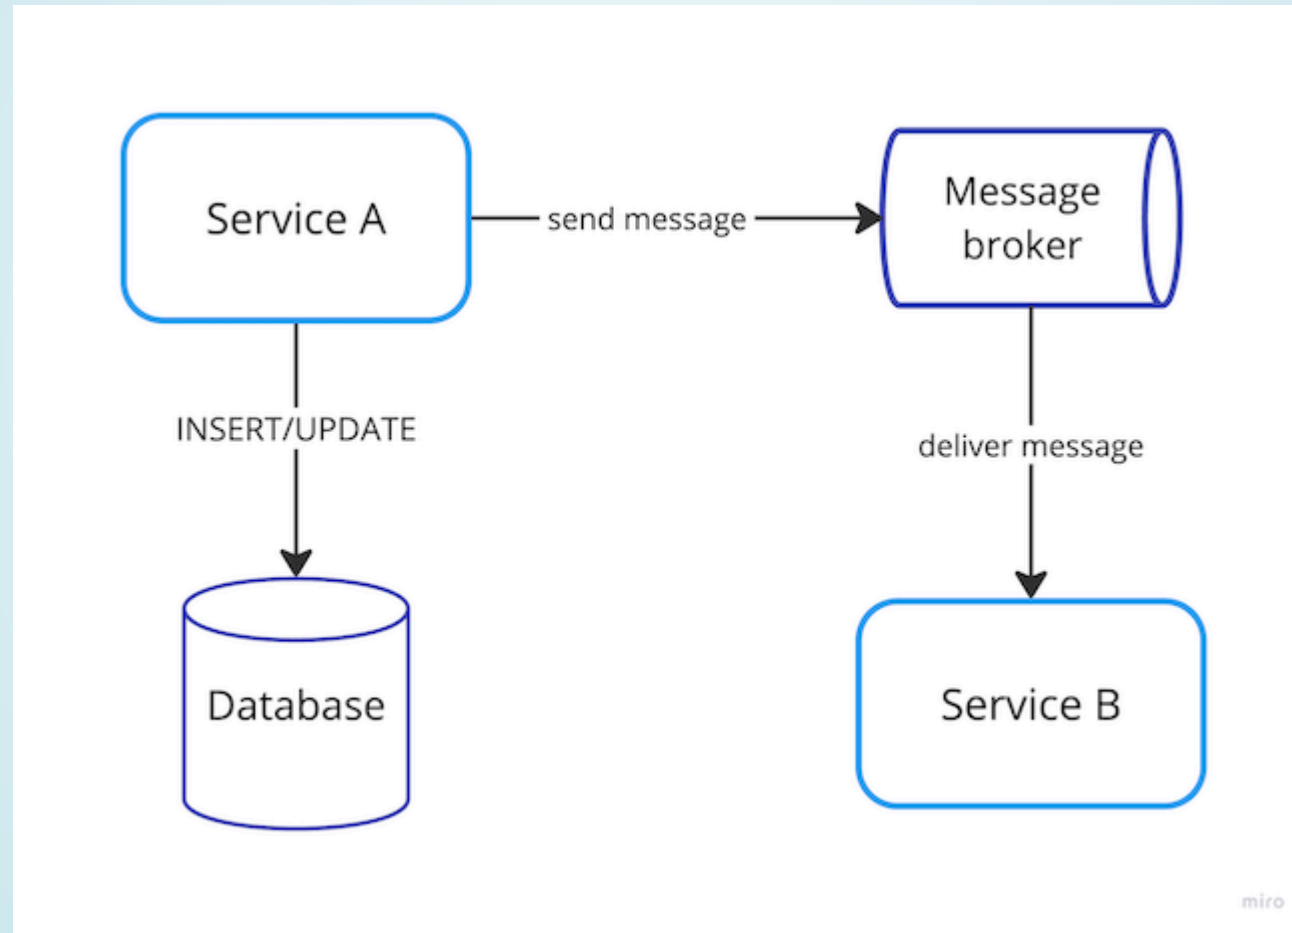
# About Zalando

Go-to-destination for fashion and lifestyle in Europe

25 countries, ~50 millions active customers

9 tech hubs, including Helsinki with ~150 employees

# Dual write problem



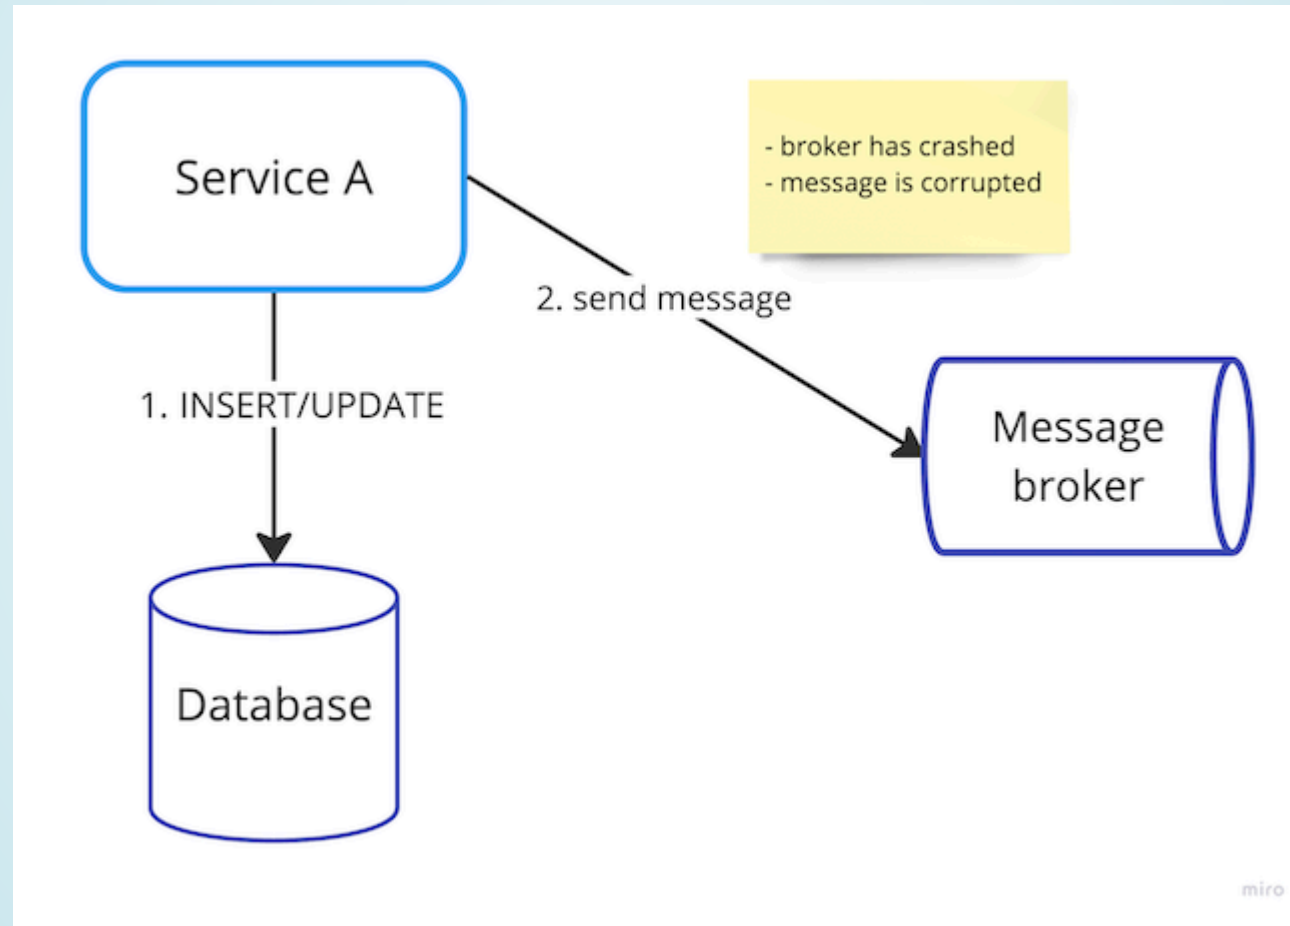
# Dual write problem

- ServiceA avoids direct sync HTTP/gRPC calls to ServiceB
- Data consistency challenge if one of writes fails
- Messages could be events notifications and state transfers
- Databases: PostgreSQL, MySQL, MariaDB, etc
- Message brokers: Kafka, SNS, NATS, etc

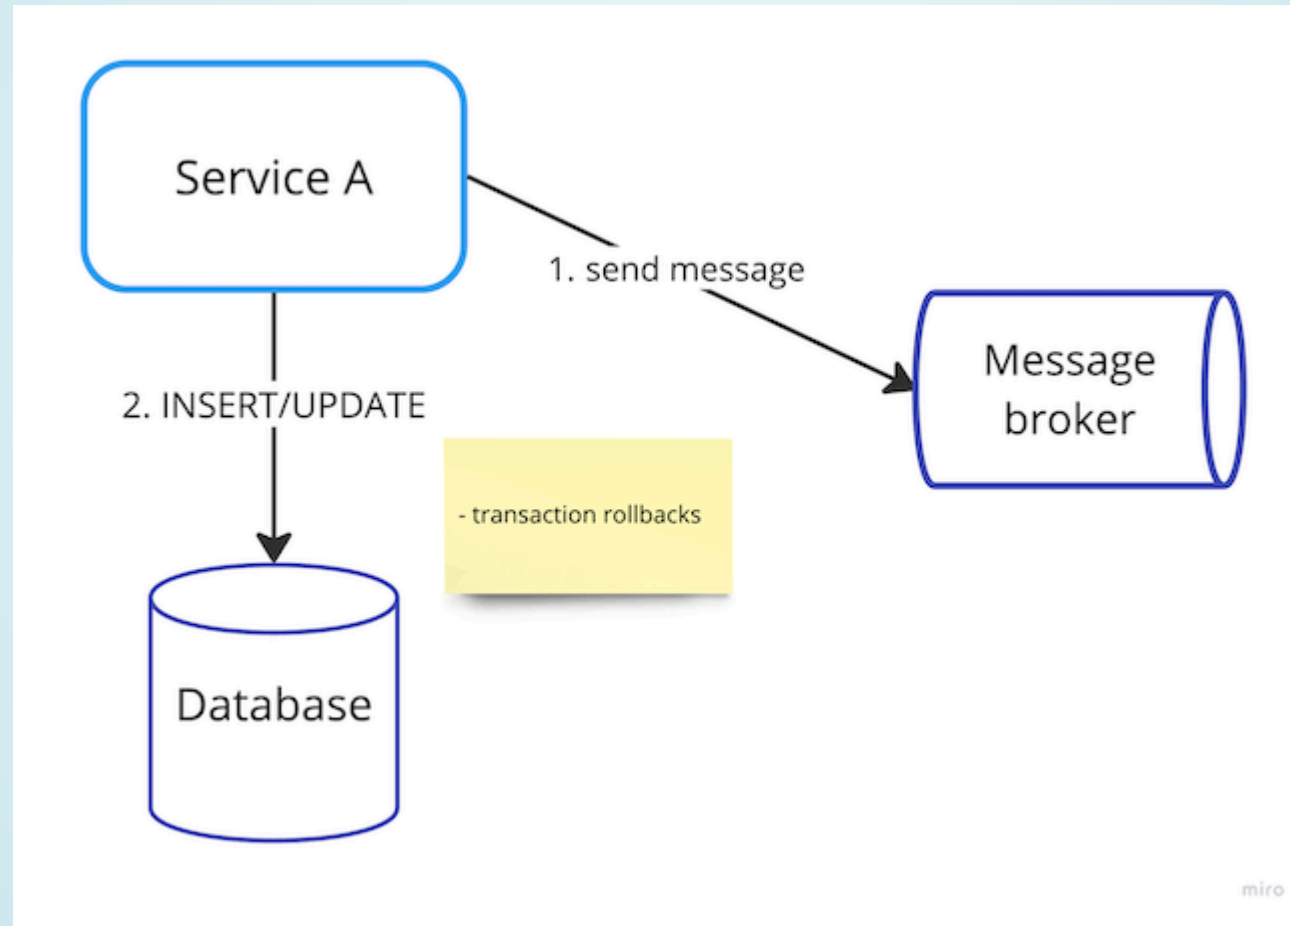
# Naive solutions

1. Write to DB first, then write to broker
2. Write to broker first, then write to DB
3. Write to broker from a DB transaction, then commit

# Database first

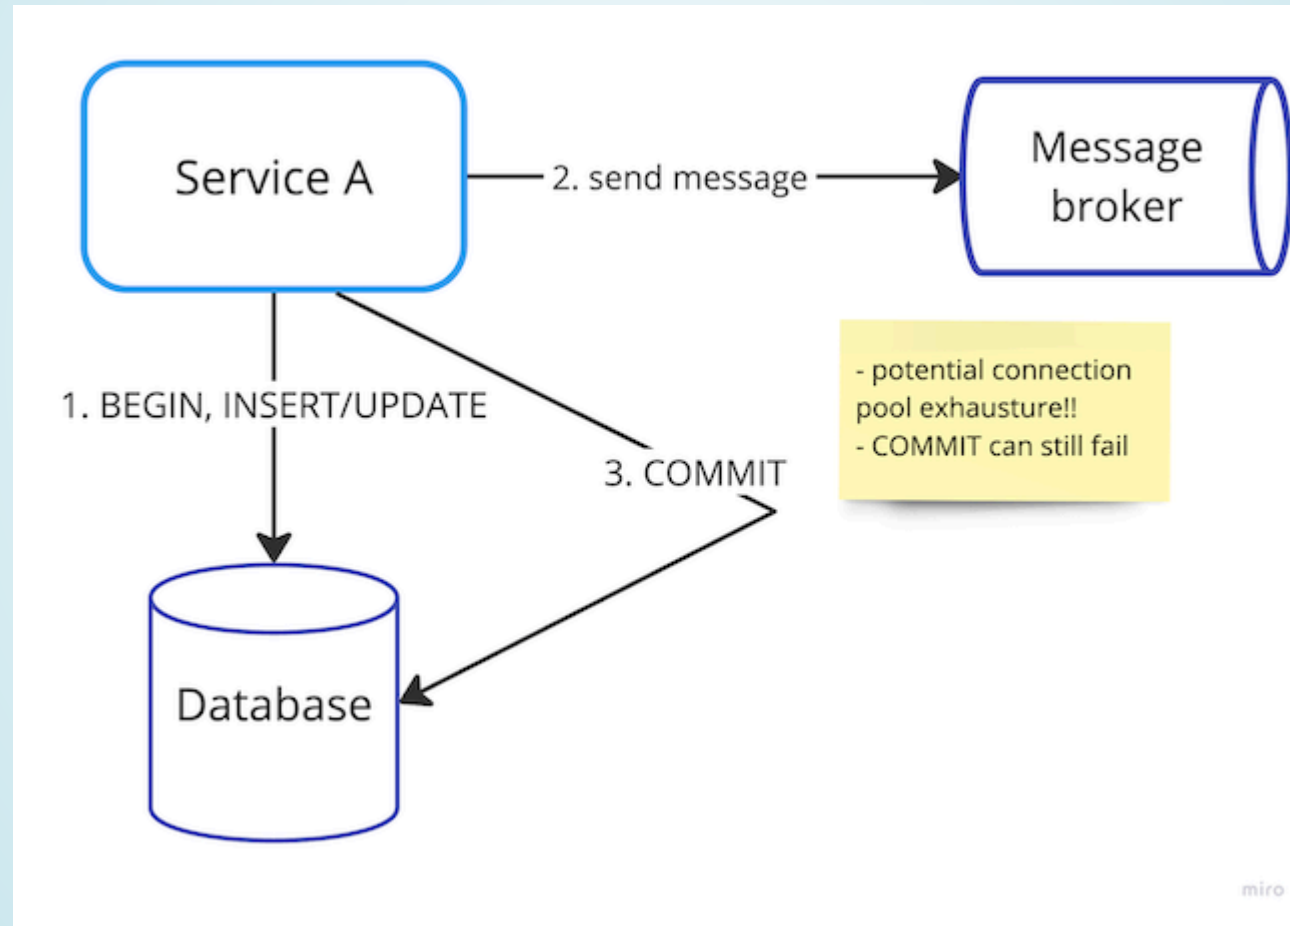


# Message broker first





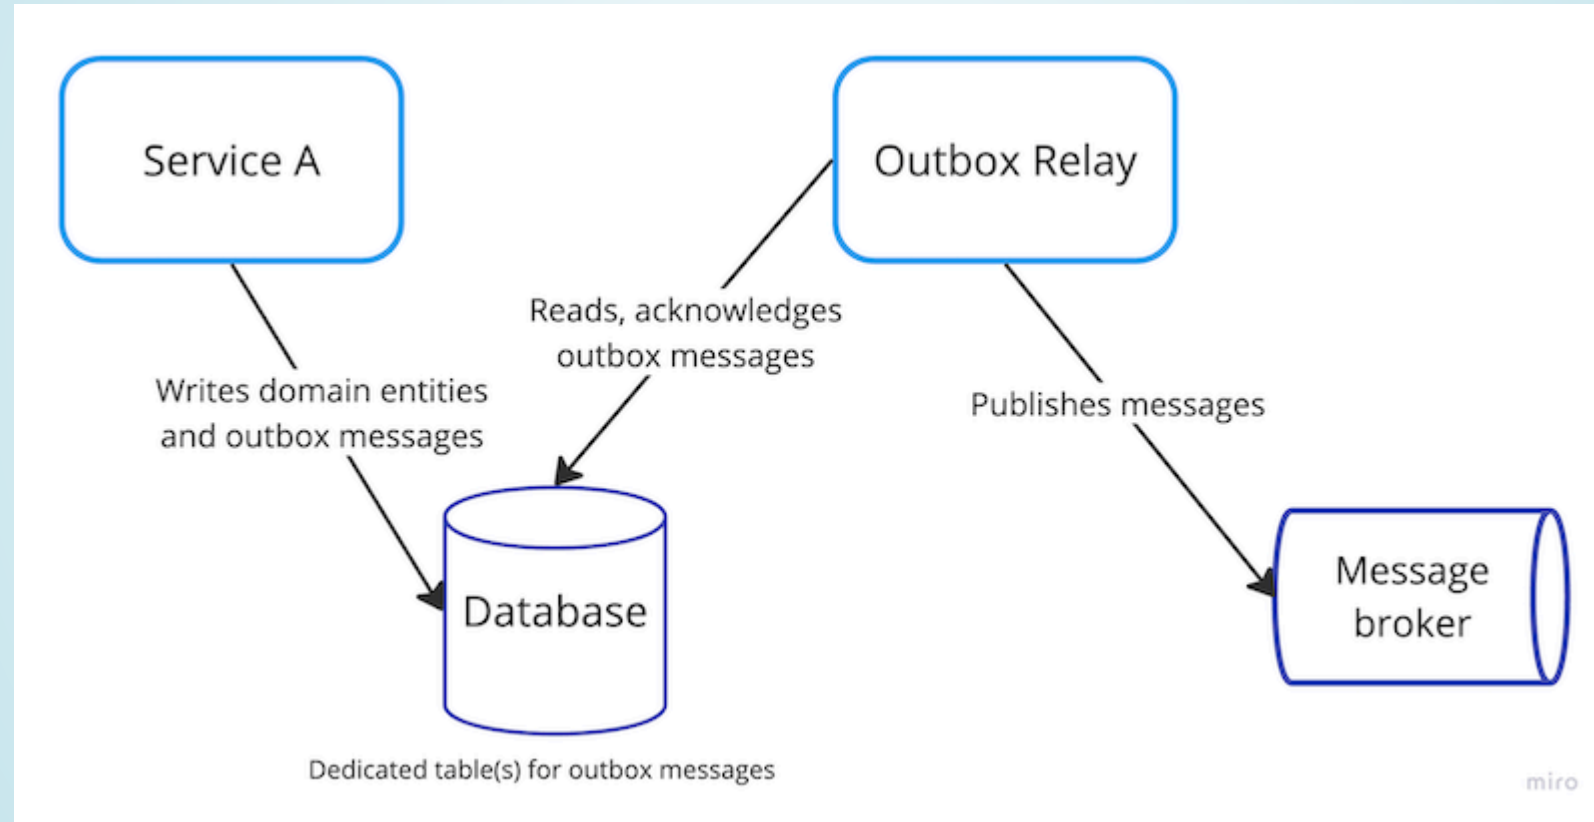
# Message from transaction



# Proper solutions

- Transactional outbox pattern
- Change data capture (CDC)
- Mix of them

# Transactional outbox pattern



# Transactional outbox pattern

- + Atomicity: domain entity vs outbox message
- + At least once delivery to message broker
- Boilerplate code, deployment unit (cronjob)
- Message consumer might need to deduplicate
- Outbox table polling introduces delay
- Autovacuum settings tuning due to MVCC

# Change data capture intro

keywords: write ahead log (WAL), logical replication, Debezium

- + Works on business and/or outbox tables
- + Near real time
- + No need for explicit SELECT/UPDATE
- Complex low-level protocol
- Not many mature products in Go ecosystem, AFAIK




Talk is cheap. Show me the code.

— *Linus Torvalds* —

AZ QUOTES

# pgx-outbox

[Go](#) [passing](#) [go report](#) [A+](#) [reference](#) [license](#) [MIT](#) [coverage](#) [80%](#) [dependencies](#) [up to date](#) [Go](#) [≥1.23](#) [Release](#) [v0.0.2](#)



## pgx-outbox

This is a simple Golang implementation for [transactional outbox](#) pattern to solve [dual writes](#) problem for PostgreSQL using [jackc/pgx](#) driver.

More advanced options are described in [Revisiting the Outbox Pattern](#) article by [Gunnar Morling](#).



# pgx driver

- for Postgres, high performance
- 11K stars at GitHub
- different interface from *database/sql*
- can be adapted to *database/sql*
- *lib/pq* is in maintenance mode



# Outbox message

```
type Message struct {  
    ID int64 // generated by Postgres  
    Broker string `validate:"required"`  
    Topic string `validate:"required"`  
    Metadata map[string]string // optional  
    Payload []byte `validate:"required,json"`  
}
```

# Outbox table

```
--table name is customizable

CREATE TABLE IF NOT EXISTS outbox_messages
(
    id                BIGINT PRIMARY KEY GENERATED ALWAYS AS IDENTITY,

    broker            TEXT                                NOT NULL,
    topic             TEXT                                NOT NULL,
    metadata          JSONB,
    payload           JSONB                                NOT NULL,

    created_at        TIMESTAMP DEFAULT CURRENT_TIMESTAMP NOT NULL,
    published_at      TIMESTAMP

);
```

# Outbox writer

```
type Writer interface {  
    // Tx is empty interface to support both pgx.Tx and *sql.Tx  
    Write(ctx, tx Tx, message Message) (int64, error)  
  
    // pgx transaction only to invoke SendBatch and Prepare methods  
    WriteBatch(ctx, tx pgx.Tx, messages []Message) ([]int64, error)  
}
```

# Add outbox writer

```
type repo struct {  
    pool *pgxpool.Pool  
  
    // new fields to use pgx-outbox  
    writer outbox.Writer  
    messageMapper ToMessageFunc[User] // can be a param instead  
}
```

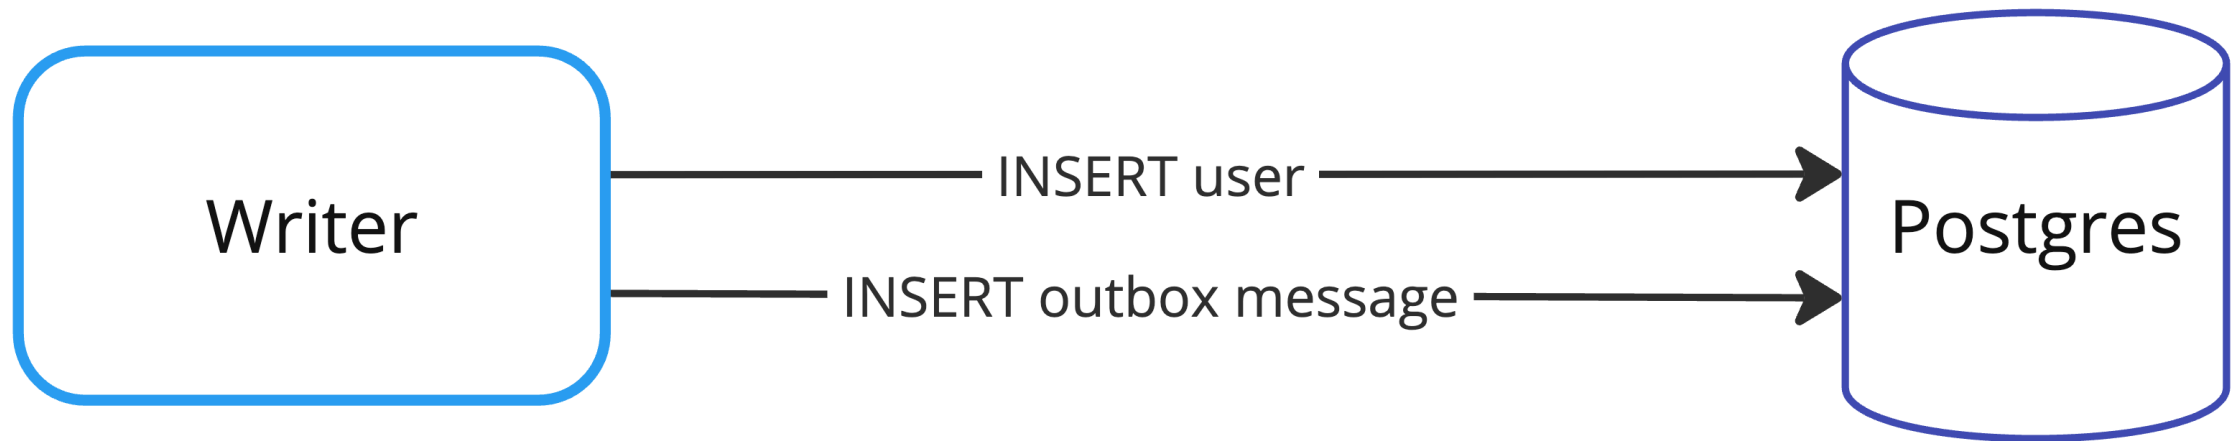
# Use outbox writer

```
func (r *repo) CreateUser(ctx, user User) (u User, txErr error) {  
    // create a transaction, commit/rollback in defer()  
  
    user, err = r.createUser(ctx, tx, user) // INSERT INTO users  
    if err != nil {  
        return u, fmt.Errorf("createUser: %w", err)  
    }  
  
    message, err := r.messageMapper(user)  
    // if err != nil {  
  
    _, err = r.writer.Write(ctx, tx, message) // INSERT INTO outbox_message  
    // if err != nil {
```

# Transaction handling

```
func (r *repo) CreateUser(ctx, user User) (u User, txErr error) {  
    tx, commitFunc, err := r.beginTx(ctx) // pool.Begin(ctx)  
    // if err != nil {  
    defer func() {  
        // commit or rollback depending on txErr  
        if cErr := commitFunc(txErr); cErr != nil {  
            txErr = fmt.Errorf("commitFunc: %w", cErr)  
        }  
    }()  
    user, err = r.createUser(ctx, tx, user)
```

# Demo outbox writer



miro

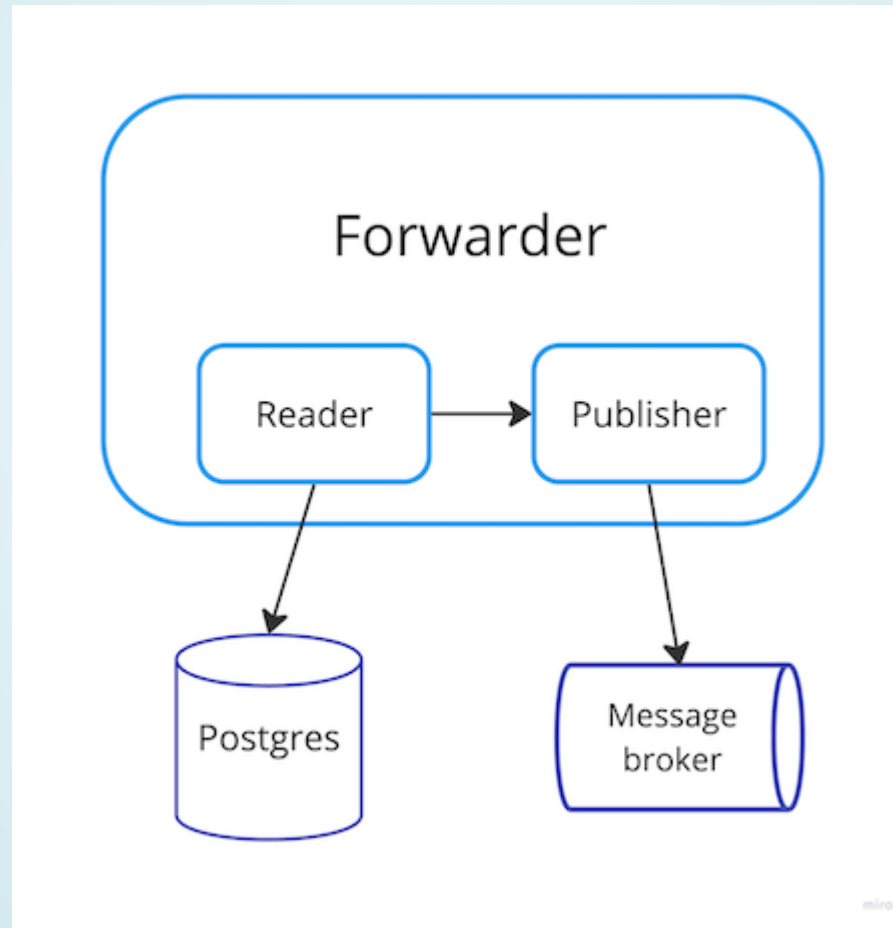
# Demo results

id	broker	topic	metadata	payload	created_at	published_at
15661	sns	topic1	{"span_id": "c49	{"id": "53e5c847-2	2025-01-14 20:32:0...	<null>
15660	sns	topic1	{"span_id": "f63	{"id": "a1ebcbc9-1	2025-01-14 20:31:5...	<null>
15659	sns	topic1	{"span_id": "df3	{"id": "587aeb48-e	2025-01-14 20:31:5...	<null>
15658	sns	topic1	{"span_id": "aa8	{"id": "1cd86f9d-9	2025-01-14 20:31:5...	<null>
15657	sns	topic1	{"span_id": "7a1	{"id": "b09279c6-8	2025-01-14 20:31:5...	<null>
15656	sns	topic1	{"span_id": "5e7	{"id": "bc1e83a9-9	2025-01-14 20:31:5...	<null>

pgx-outbox/writer user_created (11.25ms)	LOG	
pool.acquire (3.32ms)	LOG	3.32ms
query begin (1.4ms)	LOG	1.4ms
query INSERT INTO users (id, name, age) VALUES (\$1, \$2, \$3) RETURNING created_at	LOG	1.78ms
query INSERT INTO outbox_messages (broker,topic,metadata,payload) VALUES (\$1,\$2,	LOG	1.28ms
query commit (2.98ms)	LOG	2.98ms



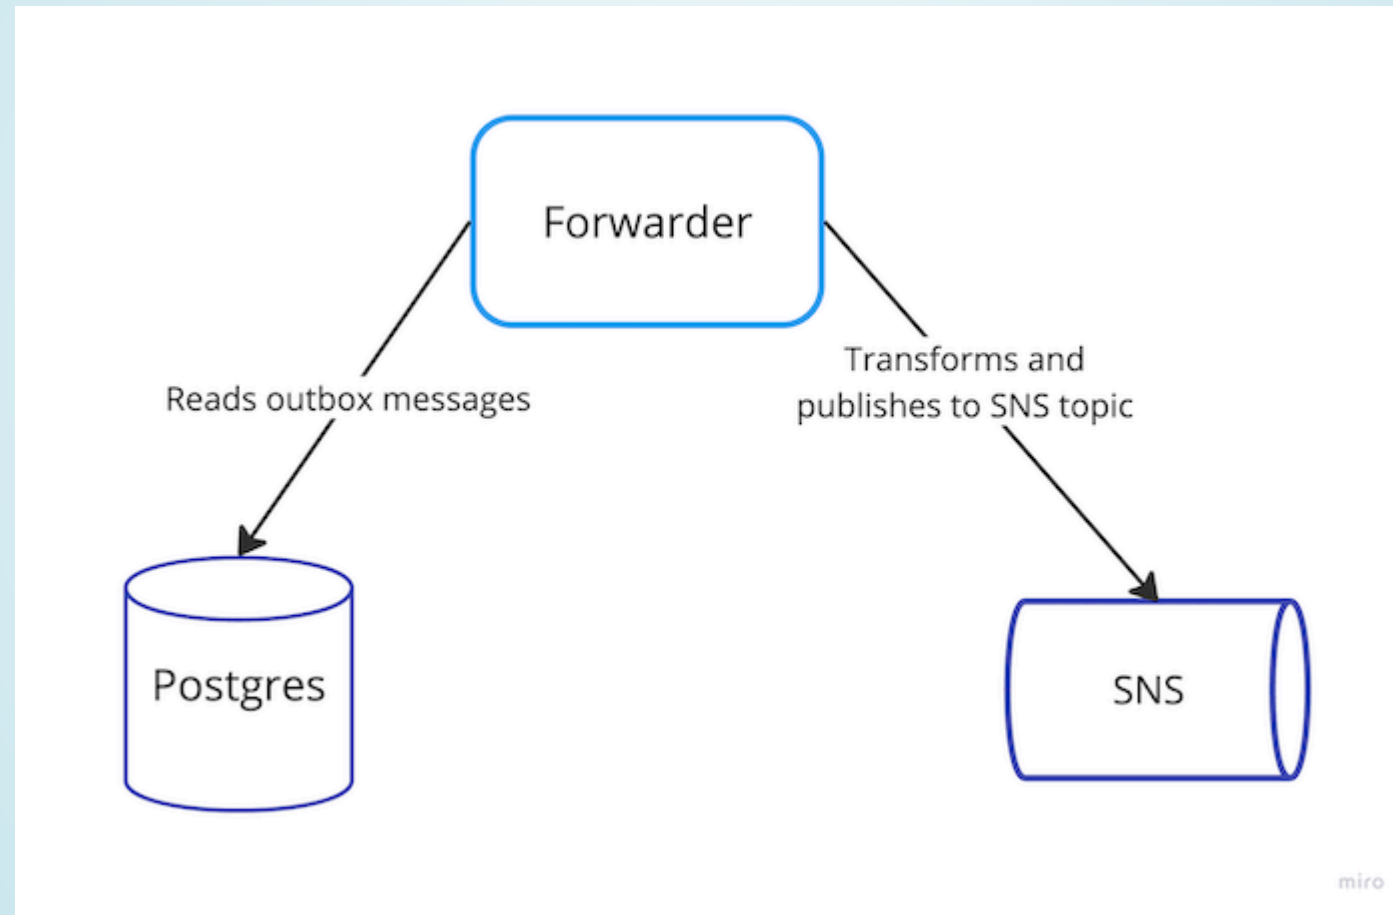
# Message relay



# Outbox forwarder

```
type Forwarder interface {  
    Forward(ctx, limit int) (ForwardOutput, error)  
}  
  
type forwarder struct {  
    reader      Reader  
    publisher   Publisher  
    filter      MessageFilter // optional  
}
```

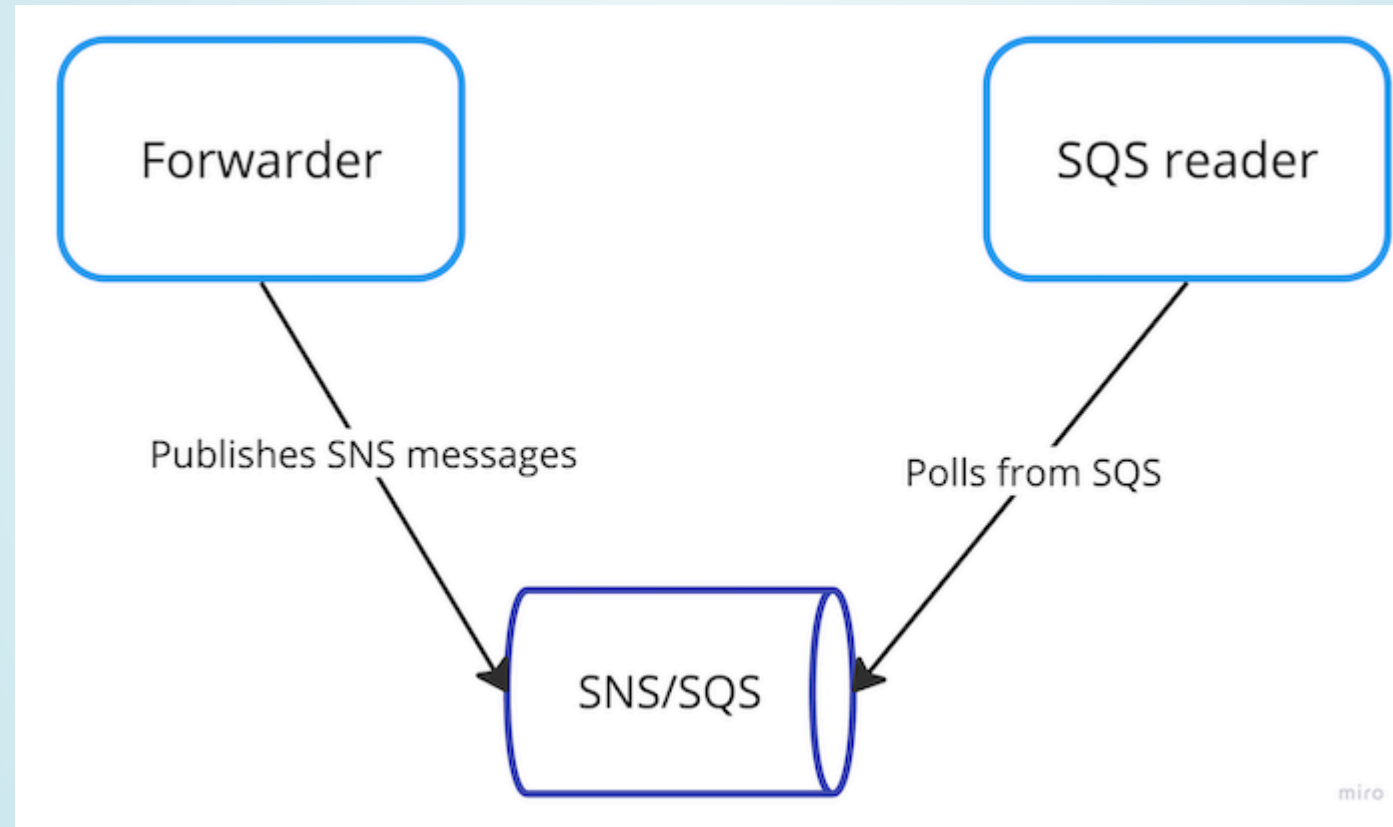
# Demo outbox forwarder



# Demo results

id	broker	topic	metadata	payload	created_at	published_at
15661	sns	topic1	{"span_id": "c49"	{"id": "53e5c847-2"	2025-01-14 20:32:0...	2025-01-14 21:36:53...
15660	sns	topic1	{"span_id": "f63"	{"id": "a1ebcbc9-1"	2025-01-14 20:31:5...	2025-01-14 21:36:53...
15659	sns	topic1	{"span_id": "df3"	{"id": "587aeb48-e"	2025-01-14 20:31:5...	2025-01-14 21:36:53...
15658	sns	topic1	{"span_id": "aa8"	{"id": "1cd86f9d-9"	2025-01-14 20:31:5...	2025-01-14 21:36:53...
15657	sns	topic1	{"span_id": "7a1"	{"id": "b09279c6-8"	2025-01-14 20:31:5...	2025-01-14 21:36:48...
15656	sns	topic1	{"span_id": "5e7"	{"id": "bc1e83a9-9"	2025-01-14 20:31:5...	2025-01-14 21:36:48...

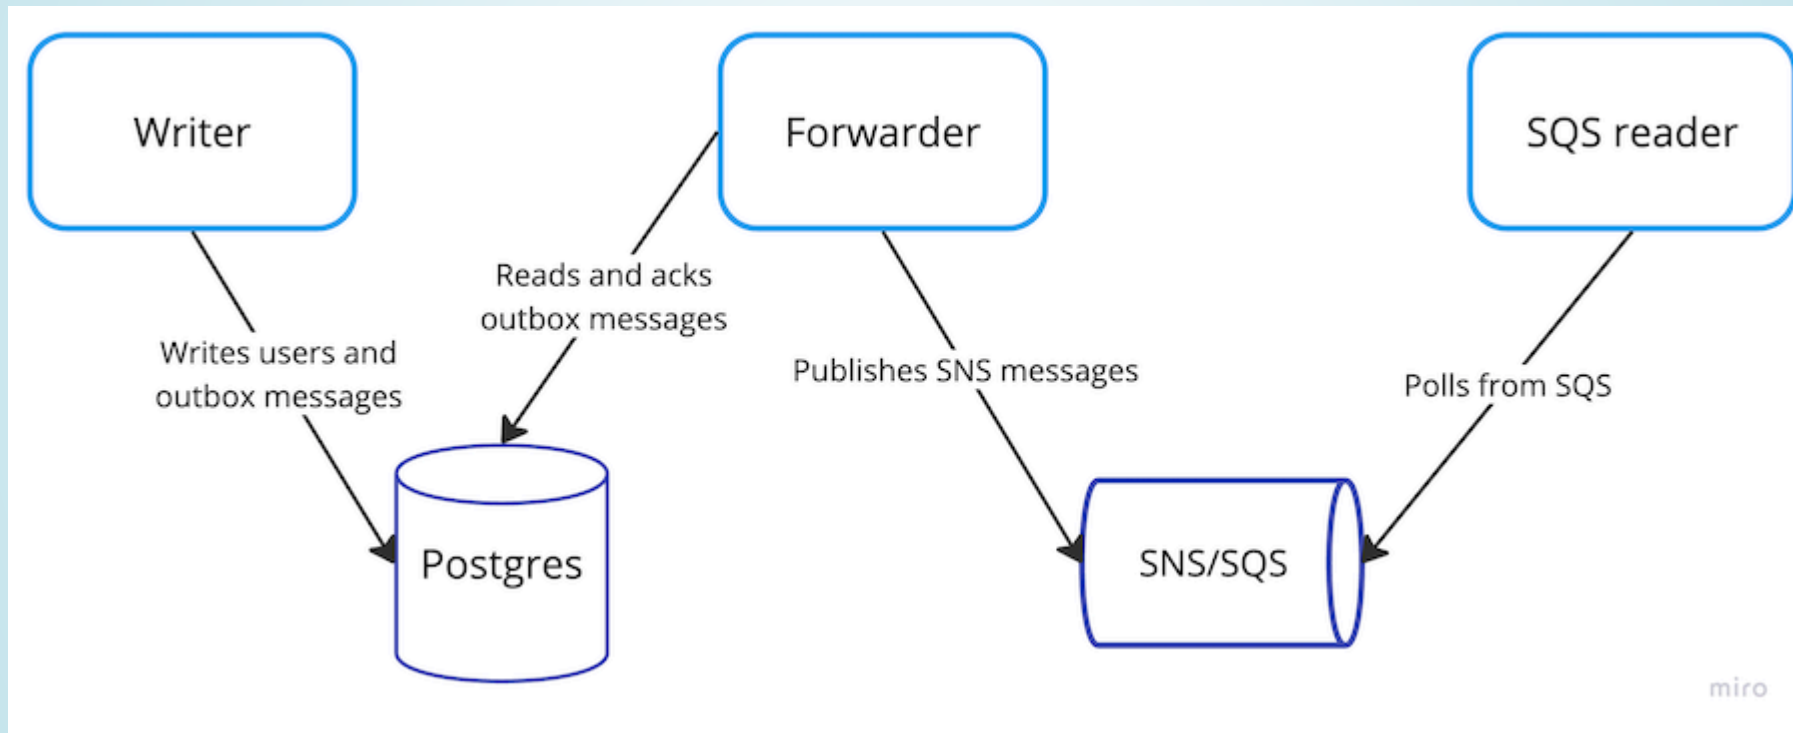
# Demo SQS reader



# Demo results

```
2025/01/21 18:35:06 Message received:
{
  "id": "fff7b31e-445c-4fb0-8748-a204dd028c2a",
  "age": 31,
  "name": "Marietta Monahan",
  "quote": "\"Beard tattooed venmo hashtag lumbersexual pickled.\" - Constance Spencer",
  "created_at": "2025-01-21T16:35:05.369315Z",
  "event_type": "user_created"
}
2025/01/21 18:35:11 Message received:
{
  "id": "e3a1688e-42da-419a-85a2-4b5e12a23533",
  "age": 60,
  "name": "Roslyn McLaughlin",
  "quote": "\"Offal post-ironic kogi carry sriracha tousled neutra s wag flannel street.\" - Demond Homenick",
  "created_at": "2025-01-21T16:35:06.884143Z",
  "event_type": "user_created"
}
```

# Demo recap



# pgx-outbox recap

Simple, generic, extensible

Writer supports *pgx* and *database/sql* types

Reader, Forwarder use *pgx*

Test coverage 80%





# Alternatives

- `watermill-sql`: SQL Pub/Sub on top of Watermill library
- `dataddo/pgq`: general queue on top of Postgres
- `jackc/pglogrepl`: Postgres logical replication, low-level
- `Trendyol/go-pq-cdc`: CDC for Postgres
- `PeerDB`: streaming from Postgres to data warehouses, queues

# Testing in pgx-outbox

Testcontainers: Postgres, LocalStack modules

Mocks: vektra/mockery

Suite: stretchr/testify/suite

Linters: testifylint

# Future plans

Add support for CDC/WAL-based Reader

Explore capabilities of PeerDB

Add support for Kafka, NATS publishers

Q & A

Thank you!

@nikolayk812

