



Introduction to ClickHouse for Gophers

Nikolai Lebedev

Helsinki Gophers Meetup | Mar. 18, 2025

About me

- Senior Software Developer
API development & Data engineering
 - Go
 - PHP
 - Python
 - ElasticSearch
 - **ClickHouse**
 - BigQuery

Problem

- High ingestion rate
- Need to aggregate data on-the-fly
- Preferably on-premises

Problem

- High ingestion rate
- Need to aggregate data on-the-fly
- Preferably on-premises

Solution:

 ClickHouse

What is ClickHouse?

- Open-source
- OLAP
- Column-oriented



Column-oriented

- Row-oriented:
 - Good for transactions
 - Slow for analytics
- Columnar:
 - Faster aggregations
 - Better compression

dt	user_id	event
1742232733	1	login
1742232733	2	login
1742232733	2	logout

dt	user_id	event
1742232733	1	login
1742232733	2	login
1742232733	2	logout

Why ClickHouse?

- Fast inserts
- Fast queries on large datasets
- Cost-effective storage with compression
- Open-source & scalable
- SQL (-like) query syntax

Drawbacks

- Not fully ACID-compliant
- Transactions are experimental
- No usual Delete/Update operations
- Might be not-so-obvious in configuration

Where can ClickHouse be used?

- Real-time analytics
- Data warehousing
- Observability
- Machine-learning, GenAI

Engines

MergeTree family	Integrations	Log family	Special
<ul style="list-style-type: none">● MergeTree● Summing● Aggregating● Collapsing● Replacing	<ul style="list-style-type: none">● S3● PostgreSQL● MySQL● Kafka● etc.	<ul style="list-style-type: none">● Log● StripeLog● TinyLog	<ul style="list-style-type: none">● Null● Distributed● URL● etc.

Engines

MergeTree family	Integrations	Log family	Special
<ul style="list-style-type: none">● MergeTree● Summing● Aggregating● Collapsing● Replacing	<ul style="list-style-type: none">● S3● PostgreSQL● MySQL● Kafka● etc.	<ul style="list-style-type: none">● Log● StripeLog● TinyLog	<ul style="list-style-type: none">● Null● Distributed● URL● etc.

Engines

MergeTree family	Integrations	Log family	Special
<ul style="list-style-type: none">● MergeTree● Summing● Aggregating● Collapsing● Replacing	<ul style="list-style-type: none">● S3● PostgreSQL● MySQL● Kafka● etc.	<ul style="list-style-type: none">● Log● StripeLog● TinyLog	<ul style="list-style-type: none">● Null● Distributed● URL● etc.

Engines

MergeTree family	Integrations	Log family	Special
<ul style="list-style-type: none">● MergeTree● Summing● Aggregating● Collapsing● Replacing	<ul style="list-style-type: none">● S3● PostgreSQL● MySQL● Kafka● etc.	<ul style="list-style-type: none">● Log● StripeLog● TinyLog	<ul style="list-style-type: none">● Null● Distributed● URL● etc.

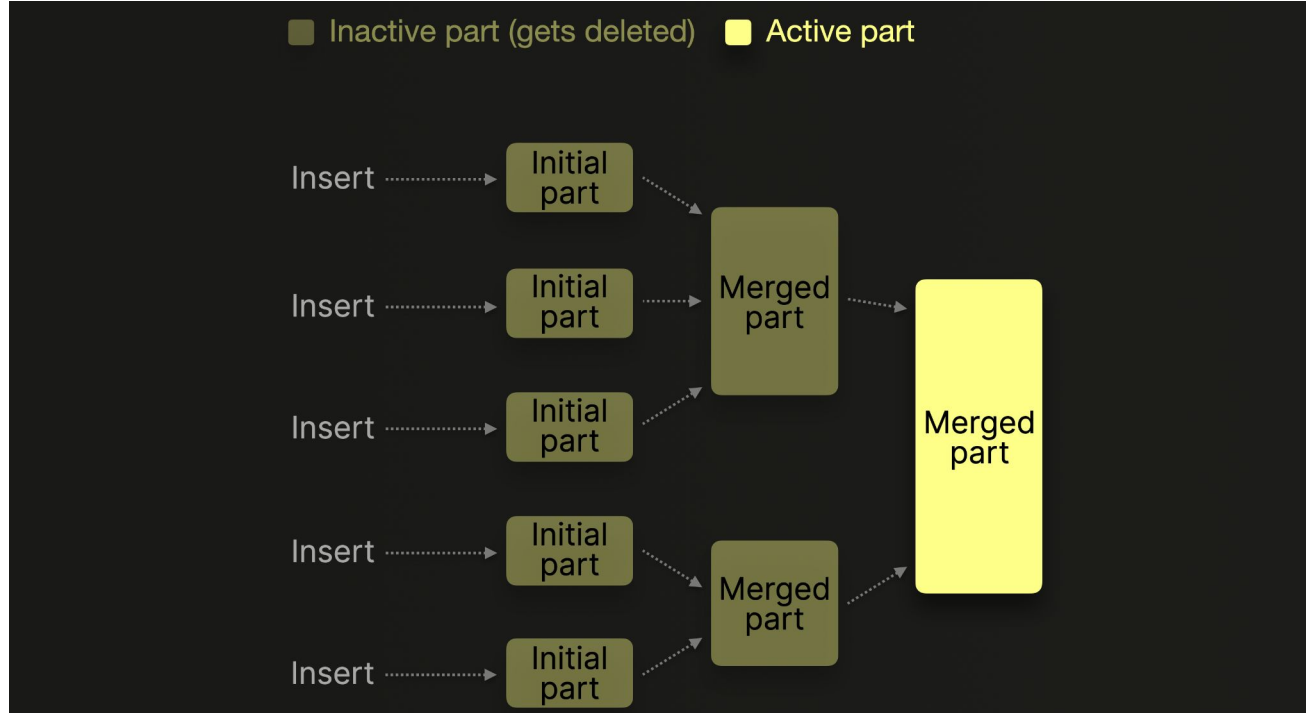
Engines

MergeTree family	Integrations	Log family	Special
<ul style="list-style-type: none">● MergeTree● Summing● Aggregating● Collapsing● Replacing	<ul style="list-style-type: none">● S3● PostgreSQL● MySQL● Kafka● etc.	<ul style="list-style-type: none">● Log● StripeLog● TinyLog	<ul style="list-style-type: none">● Null● Distributed● URL● etc.

MergeTree

```
CREATE TABLE events (  
    dt Date,  
    user_id Int32,  
    ev_type Enum('type1', 'type2', 'type3'),  
    name String,  
    INDEX ev_name (name) TYPE set(100) GRANULARITY 2,  
    PROJECTION by_event_type (SELECT * FROM events ORDER BY ev_type)  
) ENGINE = MergeTree()  
ORDER BY user_id, dt  
PRIMARY KEY user_id, dt  
PARTITION BY toYYYYMM(dt)  
TTL dt + INTERVAL 1 MONTH
```

MergeTree



Indexes: Primary key

```
CREATE TABLE events (  
  dt Date,  
  user_id Int32,  
  ev_type Enum('type1', 'type2', 'type3'),  
  name String,  
  INDEX ev_name (name) TYPE set(100) GRANULARITY 2,  
  PROJECTION by_event_type (SELECT * FROM events ORDER BY ev_type)  
) ENGINE = MergeTree()  
ORDER BY user_id, dt  
PRIMARY KEY user_id, dt  
PARTITION BY toYYYYMM(dt)  
TTL dt + INTERVAL 1 MONTH
```

Indexes: Sorting key

```
CREATE TABLE events (  
  dt Date,  
  user_id Int32,  
  ev_type Enum('type1', 'type2', 'type3'),  
  name String,  
  INDEX ev_name (name) TYPE set(100) GRANULARITY 2,  
  PROJECTION by_event_type (SELECT * FROM events ORDER BY ev_type)  
) ENGINE = MergeTree()  
ORDER BY user_id, dt  
PRIMARY KEY user_id  
PARTITION BY toYYYYMM(dt)  
TTL dt + INTERVAL 1 MONTH
```

Indexes: Projections

```
CREATE TABLE events (  
    dt Date,  
    user_id Int32,  
    ev_type Enum('type1', 'type2', 'type3'),  
    name String,  
    INDEX ev_name (name) TYPE set(100) GRANULARITY 2,  
    PROJECTION by_event_type (SELECT * FROM events ORDER BY ev_type)  
) ENGINE = MergeTree()  
ORDER BY user_id, dt  
PRIMARY KEY user_id, dt  
PARTITION BY toYYYYMM(dt)  
TTL dt + INTERVAL 1 MONTH
```


Indexes: Data skipping

```
CREATE TABLE events (  
  dt Date,  
  user_id Int32,  
  ev_type Enum('type1', 'type2', 'type3'),  
  name String,  
  INDEX ev_name (name) TYPE set(100) GRANULARITY 2,  
  PROJECTION by_event_type (SELECT * FROM events ORDER BY ev_type)  
) ENGINE = MergeTree()  
ORDER BY user_id, dt  
PRIMARY KEY user_id, dt  
PARTITION BY toYYYYMM(dt)  
TTL dt + INTERVAL 1 MONTH
```

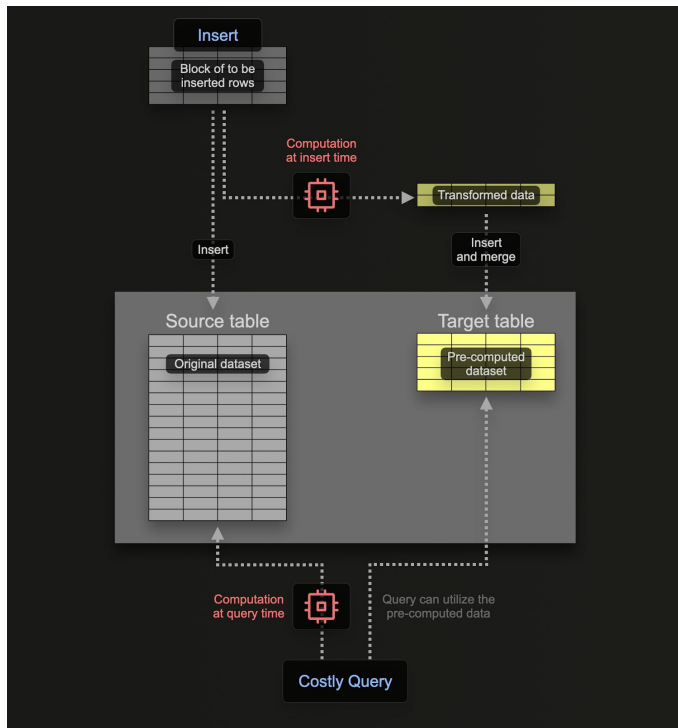
Partitions

```
CREATE TABLE events (  
  dt Date,  
  user_id Int32,  
  ev_type Enum('type1', 'type2', 'type3'),  
  name String,  
  INDEX ev_name (name) TYPE set(100) GRANULARITY 2,  
  PROJECTION by_event_type (SELECT * FROM events ORDER BY ev_type)  
) ENGINE = MergeTree()  
ORDER BY user_id, dt  
PRIMARY KEY user_id, dt  
PARTITION BY toYYYYMM(dt)  
TTL dt + INTERVAL 1 MONTH
```

Materialized Views: Refreshable



Materialized Views: Incremental



AggregatingMergeTree

-State

sumState

countState

uniqState

...

```
CREATE TABLE uniq_users (  
    dt DateTime64 NOT NULL,  
    users_uniq_cnt AggregateFunction(uniq, Nullable(Int32))  
)  
ENGINE = AggregatingMergeTree() ORDER BY (StartDate,  
CounterID);  
  
CREATE MATERIALIZED VIEW uniq_users_mv TO uniq_users  
AS SELECT  
    dt,  
    uniqState(user_id) AS users_uniq_cnt  
FROM events  
GROUP BY dt;
```


AggregatingMergeTree

-State ➔ -Merge

sumMerge

countMerge

uniqMerge

...

```
CREATE MATERIALIZED VIEW uniq_users_mv TO uniq_users
AS SELECT
    dt,
    uniqState(user_id) AS users_uniq_cnt
FROM events
GROUP BY dt;

SELECT
    dt,
    uniqMerge(users_uniq_cnt) AS users_uniq_cnt
FROM uniq_users_mv
GROUP BY dt
ORDER BY dt;
```

Go

Official clients:

- github.com/ClickHouse/ch-go
- github.com/ClickHouse/clickhouse-go

Alternative clients:

- github.com/vahid-sohrabloo/chconn
- github.com/uptrace/go-clickhouse

ch-go

- Column-oriented
- Suitable for heavy inserts
- No reflection or interface{}
- More efficient
- No pooling (available via chpool)
- Not goroutine-safe
- Not database/sql compatible

clickhouse-go

- More high-level (struct tags, helper functions, pooling, etc.)
- Suitable for queries
- *database/sql* -compatible

github.com/ClickHouse/ch-bench

Name	Time	RAM	Ratio
ClickHouse/ch-go (Go)	401ms	9M	~1x
clickhouse-client (C++)	387ms	91M	~1x
vahid-sohrabloo/chconn (Go)	472ms	9M	~1x
clickhouse-cpp (C++)	516ms	6.9M	1.47x
clickhouse_driver (Rust)	614ms	9M	1.72x
curl (C, HTTP)	3.7s	10M	9x
clickhouse-client (Java, HTTP)	6.4s	121M	16x
clickhouse-jdbc (Java, HTTP)	7.2s	120M	18x
loyd/clickhouse.rs (Rust, HTTP)	10s	7.2M	28x
uptrace (Go) ^[1]	22s	13M	55x
clickhouse-driver (Python)	37s	60M	106x
ClickHouse/clickhouse-go (Go) ^[1]	46.8s	23M	117x
mailru/go-clickhouse (Go, HTTP)	4m13s	13M	729x

Struct tags

```
type Event struct {  
    Dt          time.Time `ch:"dt"`  
    UserID      string    `ch:"user_id"`  
    EType       string    `ch:"ev_type"`  
    Name        string    `ch:"name"`  
}  
  
var result []Event  
  
err := conn.Select(ctx, &result, "SELECT dt, user_id, ev_type, name FROM events")  
...
```

Batch insert: Clickhouse API

```
batch, err := conn.PrepareBatch(ctx, "INSERT INTO events")
...
for i := 0; i < 1000; i++ {
    err := batch.Append(v1,v2,v3,v4)
    ...
}
batch.Send()
```

Batch insert: database/sql

```
tx, err := conn.Begin()
...
tx, err := scope.Prepare("INSERT INTO events")
...
for i := 0; i < 1000; i++ {
    _, err := batch.Exec(v1,v2,v3,v4)
    ...
}
tx.Commit()
```


Async insert: Clickhouse API

```
err := conn.AsyncInsert(ctx, `INSERT INTO events VALUES (?, ?, ?, ?)\`, ...);
```

Async insert: database/sql

```
// false - wait_for_async_inserts=false
ctx := clickhouse.Context(context.Background(), clickhouse.WithStdAsync(false))
for i := 0; i < 100; i++ {
    _, err := conn.ExecContext(ctx, `INSERT INTO events VALUES (?, ?, ?, ?)`, ...)
    ...
}
```

Columnar insert

```
batch, err := conn.PrepareBatch(context.Background(), "INSERT INTO events")
...
var (
    tsCol []time.Time
    ...
)
for i := 0; i < 1_000; i++ {
    col1 = append(col1, time.Now())
    ...
}
err := batch.Column(0).Append(col1)
...
```

Context

```
ctx, cancel := context.WithCancel(context.Background())  
  
...  
  
cancel()  
  
...
```

```
ctx, cancel = context.WithDeadline(  
    context.Background(),  
    time.Now().Add(time.Second * 10)  
)  
  
defer cancel()
```

Context

```
ctx = clickhouse.Context(  
    context.Background(),  
    clickhouse.WithQueryID(queryId.String())  
)
```

Context

```
ctx := clickhouse.Context(  
    context.Background(),  
    clickhouse.WithSettings(clickhouse.Settings{  
        "allow_experimental_object_type": "1",  
    }))
```

Thank you for your attention!

LinkedIn: <https://linkedin.com/in/nikolai-lebedev/>